

Satellite Image Segmentation By Land Use

By Luan Nguyen, Nick(Chi Yen Yu), Juan, Ethan, Alan

A dark blue diagonal gradient bar that starts from the bottom left and extends towards the top right, covering the lower half of the slide.

Background

Takes satellite images taken by Sentinel-2 categorize use of land into 10 different categories.

Examples of practical use:

City planning / Real Estate

Environmental Monitoring




Dataset

Satellite images taken by Sentinel-2.

Containing 27,000 images (64x64, 10m definition).











Divided into 10 categories of usage.

Two options: RGB or 12-band.



About this directory

This directory only contains JPG images with Red, Green and Blue bands of images obtained from the Sentinel-2 satellite.

 AnnualCrop 3000 files	 Forest 3000 files	 HerbaceousVegetation 3000 files	 Highway 2500 files	 Industrial 2500 files
 Pasture 2000 files	 PermanentCrop 2500 files	 Residential 3000 files	 River 2500 files	 SeaLake 3000 files


Method

Classification: Using a two layer CNN model with a final linear layer to predict the class of the image.

Segmentation: Using a U-Net network to predict the class of each pixel in the image.

Each image has a category. We first create a label image for each category and each input image. The corresponding category will have all 1.0, the other will have all 0.0.

We then apply the model to the input image, which will give us the predicted probabilities for each pixel and each category. We then calculate the loss function. Before backpropagating, we apply a gaussian filter to the loss.



Evaluation

Since the test set only has the categorical label for the whole image, we have to use an averaging algorithm to get the predicted label.

Available algorithms:

- Average all category then select max.
- Select max for each pixel then average.
- Fully connected neural network (freeze segmentation network).

Metrics: accuracy, precision and recall.

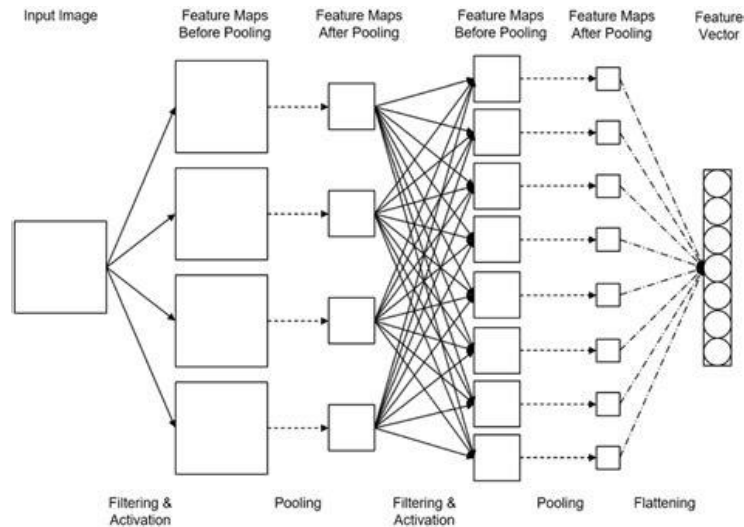
Results – Part I

Model: (Classifier) 2 layers convolution network + dense layer.

Accuracy: 68%.

TABLE II: Classification accuracy (%) of different training-test splits on the EuroSAT dataset.

Method	10/90	20/80	30/70	40/60	50/50	60/40	70/30	80/20	90/10
BoVW (SVM, SIFT, k = 10)	54.54	56.13	56.77	57.06	57.22	57.47	57.71	58.55	58.44
BoVW (SVM, SIFT, k = 100)	63.07	64.80	65.50	66.16	66.25	66.34	66.50	67.22	66.18
BoVW (SVM, SIFT, k = 500)	65.62	67.26	68.01	68.52	68.61	68.74	69.07	70.05	69.54
CNN (two layers)	75.88	79.84	81.29	83.04	84.48	85.77	87.24	87.96	88.66
ResNet-50	75.06	88.53	93.75	94.01	94.45	95.26	95.32	96.43	96.37
GoogleNet	77.37	90.97	90.57	91.62	94.96	95.54	95.70	96.02	96.17



Results – Part II


Model: U-Net Architecture.

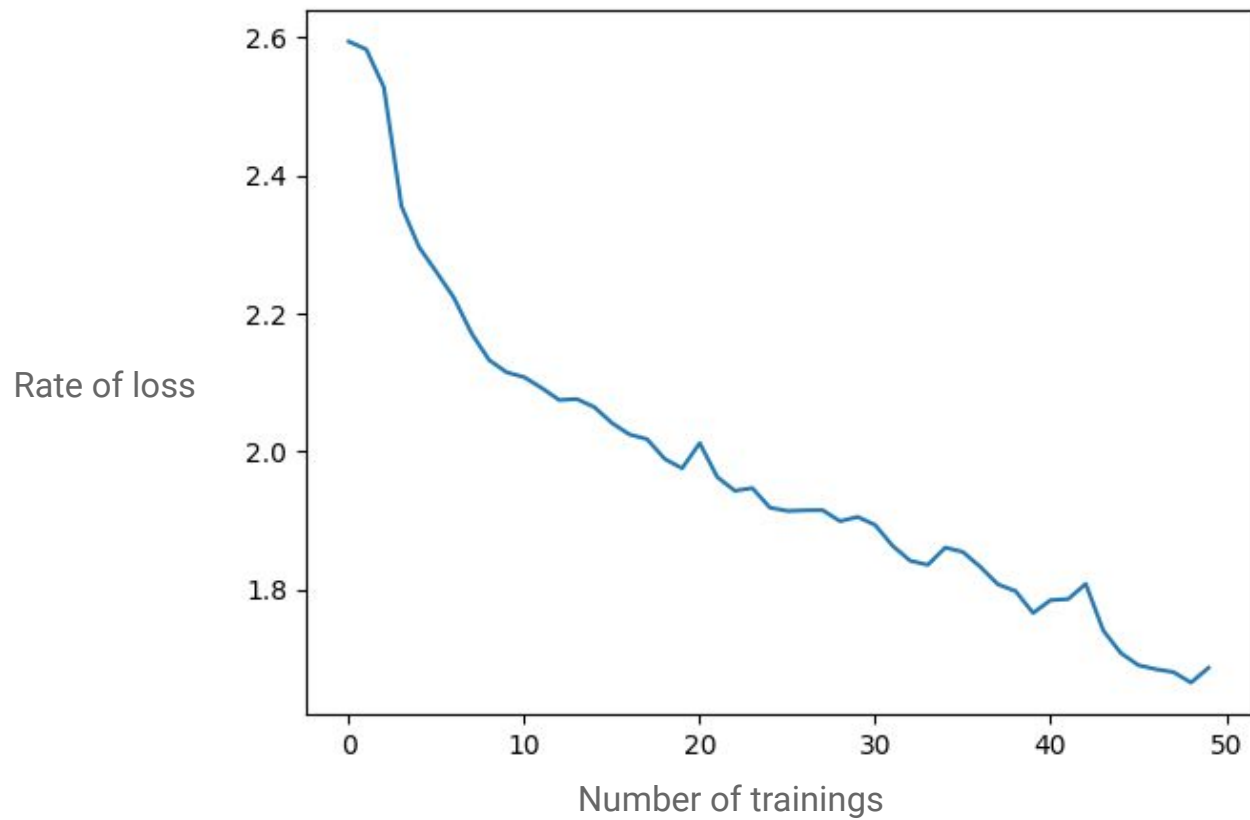
($8 \rightarrow 12 \rightarrow 16 \rightarrow 20 \rightarrow 20+16 \rightarrow 16+12 \rightarrow 12+8 \rightarrow 8$
 $\rightarrow 10$ classes).

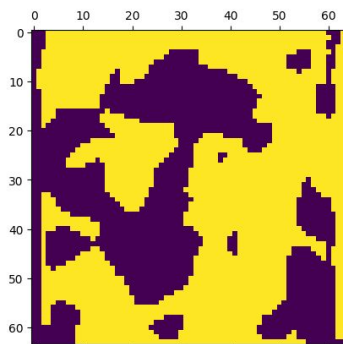

Accuracy: 86%.

Precision: 93%.


Recall: 91%.









Annual Crop




Herbaceous
Vegetation




Highway




Industrial



Permanent Crop



Residential



River

```

1 # %x
2
3 import torch
4 import torch.utils.data
5 import torch.optim
6 import torchvision as vision
7 import torchvision.transforms.v2 as v2
8 import tqdm
9 import matplotlib.pyplot as pyplot
10
11 import models
12
13 def gaussian_filter(kernel_size, sigma=1, muu=0):
14     import numpy as np
15     # Initializing value of x,y as grid of kernel size
16     # in the range of kernel size
17     x, y = np.meshgrid(np.linspace(-1, 1, kernel_size),
18                       np.linspace(-1, 1, kernel_size))
19     dst = np.sqrt(x**2+y**2)
20     # Calculating Gaussian filter
21     gauss = np.exp(-((dst-muu)**2 / (2.0 * sigma**2)))
22     return gauss
23
24 Run Cell | Run Above | Debug Cell
25 # %x
26
27 BATCH_SIZE = 2000
28 TRAIN_SET_SIZE = 0.8
29
30 def target_transform(label: int) -> torch.Tensor:
31     """Transform label into segmentation matrix."""
32     l = torch.zeros((10, 64, 64))
33     l[label] = torch.ones((64, 64))
34     return l
35
36 def target_transform_2(label: int) -> torch.Tensor:
37     """Transform label into one-hot encoded."""
38     l = torch.zeros(10)
39     l[label] = 1
40     return l
41
42
43 dataset = vision.datasets.ImageFolder(
44     root="dataset",
45     transform=vision.transforms.Compose([
46         v2.PILToTensor(),
47         v2.ToDtype(torch.float32, scale=True),
48     ]),
49     target_transform=target_transform)
50
51 train_set, test_set = torch.utils.data.random_split(dataset, [TRAIN_SET_SIZE, 1 - TRAIN_SET_SIZE])
52 train_loader = torch.utils.data.DataLoader(dataset=train_set, batch_size=BATCH_SIZE, shuffle=True)
53 test_loader = torch.utils.data.DataLoader(dataset=test_set, batch_size=BATCH_SIZE, shuffle=False)
54
55 print(f"The dataset has {len(dataset)} samples.")
56 print(f"The train set has {len(train_set)} samples.")
57 print(f"The test set has {len(test_set)} samples.")
58

```

```

60 unet = models.UNet(nc_classes=10).cuda()
61
62 # criterion = torch.nn.BCELoss(reduction="mean")
63 criterion2 = torch.nn.BCELoss(reduction="none")
64 # optimizer = torch.optim.Adam(cls.parameters(), lr=1e-3)
65 optimizer = torch.optim.Adam(unet.parameters(), lr=1e-3)
66
67 loss_filter = torch.Tensor(gaussian_filter(64)).cuda()
68
69 loss_hist = []
70
71 Run Cell | Run Above | Debug Cell
72 # %x
73
74 EPOCHS = 50
75
76 for epoch in range(EPOCHS):
77     total_loss = 0.0
78     for inputs, labels in tqdm.tqdm(train_loader):
79         inputs = inputs.cuda()
80         labels = labels.cuda()
81         # y = cls(inputs)
82         y = unet(inputs)
83         loss = criterion2(y, labels)
84         loss = (loss * loss_filter).mean()
85         optimizer.zero_grad()
86         loss.backward()
87         optimizer.step()
88         # loss_hist.append(loss.item())
89         total_loss += loss.item()
90     loss_hist.append(total_loss)
91
92 pyplot.plot(loss_hist)
93
94 tp = 0
95 tn = 0
96 fp = 0
97 fn = 0
98
99 for inputs, labels in tqdm.tqdm(test_loader):
100     inputs = inputs.cuda()
101     labels = labels.cuda()
102     # y = cls(inputs)
103     y = unet(inputs)
104     # acc = torch.argmax(y, dim=1) == torch.argmax(labels, dim=1)
105     # accuracy += acc.sum()
106     y = torch.argmax(y, dim=1).to(bool)
107     labels = torch.argmax(labels, dim=1).to(bool)
108     tp += ((y == True) & (labels == True)).sum()
109     tn += ((y == False) & (labels == False)).sum()
110     fp += ((y == True) & (labels == False)).sum()
111     fn += ((y == False) & (labels == True)).sum()
112
113 accuracy = (tp + tn) / (tp + tn + fp + fn)
114 precision = tp / (tp + fp)
115 recall = tp / (tp + fn)
116 print(f"Accuracy = {accuracy}")
117 print(f"Precision = {precision}")
118 print(f"Recall = {recall}")
119
120 Run Cell | Run Above | Debug Cell
121 # %x
122
123 torch.save(
124     {
125         "model_state_dict": unet.state_dict(),
126         "optimizer_state_dict": optimizer.state_dict(),
127     },
128     "outputs/state.pt")
129
130
131

```